

XIANFENG TANG

[Homepage](#) ◊ [Email](#) ◊ [LinkedIn](#) ◊ [Google Scholar](#)

INTRODUCTION

Xianfeng Tang is an experienced researcher with five years of hands-on work in large foundation models and their real-world applications, including Amazon Rufus and Amazon Ads foundation models. His research spans a broad range of topics such as supervised fine-tuning, reinforcement learning, reasoning, agentic capabilities, multi-modality, knowledge editing and unlearning, trust and safety, retrieval augmentation, and e-commerce applications.

RECENT EMPLOYMENT

Senior Applied Scientist , founding member of the Amazon ADs foundation model team	03/2025 - Present
Senior Applied Scientist , built Rufus , Amazon's Shopping AI	02/2023 - 02/2025
Applied Scientist , improved Amazon search experiences	01/2021 - 01/2023

EDUCATION

Ph.D. in Computer Science , Penn State University Improve the robustness, fairness, and explainability of graph neural network models.	08/2016 - 12/2020
B.S in Computer Science , University of Science and Technology of China Excellent Graduate Student (honored class)	08/2012 - 06/2016

PUBLICATIONS

See the complete publication list on [Google Scholar](#).

Post-training

- *Mitigating Heterogeneous Token Overfitting in LLM Knowledge Editing*, [PDF](#)
- *SFT or RL? An Early Investigation into Training R1-Like Reasoning Large Vision-Language Models*, [PDF](#)
- *Keeping an Eye on LLM Unlearning: The Hidden Risk and Remedy*, [PDF](#)
- *A General Framework to Enhance Fine-Tuning-Based LLM Unlearning*, [PDF](#)
- *Does your LLM truly unlearn? An embarrassingly simple approach to recover unlearned knowledge*, [PDF](#)
- *Divide-Verify-Refine: Aligning LLM Responses with Complex Instructions*, [PDF](#)

Reasoning

- *Efficient Long CoT Reasoning in Small Language Models*, [PDF](#)
- *Stepwise Perplexity-Guided Refinement for Efficient Chain-of-Thought Reasoning in Large Language Models*, [PDF](#)
- *A Theoretical Understanding of Chain-of-Thought: Coherent Reasoning and Error-Aware Demonstration*, [PDF](#)
- *Graph Chain-of-Thought: Augmenting Large Language Models by Reasoning on Graphs*, [PDF](#)
- *m1: Unleash the Potential of Test-Time Scaling for Medical Reasoning with Large Language Models*, [PDF](#)
- *Reasoning with Graphs: Structuring Implicit Knowledge to Enhance LLMs Reasoning*, [PDF](#)

Agentic Capability

- *Attention Knows Whom to Trust: Attention-Based Trust Management for LLM Multi-Agent Systems*, [PDF](#)
- *Comprehensive Vulnerability Analysis is Necessary for Trustworthy LLM-MAS*, [PDF](#)
- *RRO: LLM Agent Optimization Through Rising Reward Trajectories*, [PDF](#)

- *Cite Before You Speak: Enhancing Context-Response Grounding in E-Commerce Conversational LLM-Agents*, [PDF](#)

Efficiency

- *Examples as the Prompt: A Scalable Approach for Efficient LLM Adaptation in E-Commerce*, [PDF](#)
- *Resmoe: Space-efficient compression of mixture of experts LLMs via residual restoration*, [PDF](#)
- *Unlocking Efficient, Scalable, and Continual Knowledge Editing with Basis-Level Representation Fine-Tuning*, [PDF](#)
- *Blendfilter: Advancing retrieval-augmented large language models via query generation blending and knowledge filtering*, [PDF](#)
- *Learning with Less: Knowledge Distillation from Large Language Models via Unlabeled Data*, [PDF](#)
- *Simrag: Self-improving retrieval-augmented generation for adapting large language models to specialized domains*, [PDF](#)
- *Harnessing the Unseen: The Hidden Influence of Intrinsic Knowledge in Long-Context Language Models*, [PDF](#)

Applications

- *Exploring Query Understanding for Amazon Product Search*, [PDF](#)
- *Towards Unified Multi-modal Personalization: Large Vision-Language Models for Generative Recommendation and Beyond*, [PDF](#)
- *Sequential LLM Framework for Fashion Recommendation*, [PDF](#)
- *IHEval: Evaluating Language Models on Following the Instruction Hierarchy*, [PDF](#)
- *Shopping MMLU: A Massive Multi-task Online Shopping Benchmark for Large Language Models*, [PDF](#)
- *Large Language Models in the Clinic: A Comprehensive Benchmark*, [PDF](#)
- *Amazon-M2: A Multilingual Multi-locale Shopping Session Dataset for Recommendation and Text Generation*, NeurIPS 2023 [PDF](#)

HONORS AND AWARDS

- Travel Awards from AAAI, KDD, CIKM, etc. 2019 - 2020
- Rednor Graduate Fellowship, Penn State 2016
- 3rd Prize of Student Cluster Competition, SC14 2014
- First Prize, Olympiad in Mathematics, Anhui Province, China 2011